

BGP Divergence Continued

Computer Networks
Dr. Jorge A. Cobb



Statically Solving an SPP

- Given an SPP, can we check if it is solvable?
- Sure,
 - a) enumerate all possible states
 - b) For each state, check if it is stable
- Checking for safety is even worse (check all sub-instances)
- Solvability of SPP is NP-Complete
 - a) Exponential complexity!! (that we know of)

A Sufficient Condition for Sanity

- If an instance of SPP has no **dispute wheel**, then

Static (SPP)

solvable

unique solution

**all sub-problems
uniquely solvable**


Dynamic SPP

converges

predictable restoration

**robust with respect to
link/node failures (safe)**


Dispute Wheels




- u_0, u_1, \dots, u_k are nodes (not necessarily distinct)
- R_i is a path from u_i to $u_{(i+1)}$
- Q_i is a path from u_i to 0
- Q_i and $R_i Q_{(i+1)} \in \mathbf{P}^{u_i}$
- $\lambda^{u_i}(Q_i) < \lambda^{u_i}(R_i Q_{(i+1)})$

Example of Dispute Wheel


“Bad Triangle”




Its “dispute wheel”



“Bad Gadget”



A “dispute wheel”





Sufficient Condition for a Solution

- **Theorem:** if an SPP instance does not have a dispute wheel, then it has a solution
- Note: this is a sufficient but not necessary condition
 - a) A dispute wheel could exist and still we have a solution
(can you add nodes to bad triangle and make it converge while still having a dispute wheel???)
- We will build a spanning tree such that:
 - a) If we complete the construction then there is a solution
 - b) If we “get stuck”, then there is a dispute wheel
 - Equivalently (contrapositive):
no dispute wheel → spanning tree built → solution exists



Ensuring solution IS a spanning tree

- Some solutions have some nodes with an empty path (not a spanning tree).
- We modify the SPP instance by adding an additional node x , and adding for each node u , the path $(u \ x \ 0)$, that is ranked lowest among all paths at u .
- You can show the solutions are the same as before, except no one has an empty path.



Tree construction

- For a tree T (not necessarily spanning), $V(T)$ are its vertices (nodes)
- Let S be our SPP instance.
- At each step of our construction, the sub-instance S' obtained from S by removing all nodes not in $V(T)$ and all paths with nodes not in $V(T)$ **is stable**.



T is stable

Note that $(1, 0)$ remains since both 1 and 0 are in T


Adding a node to the tree

- A path P is said to be *consistent* with tree T if once P encounters a node in T , the rest of P is along T
 - I.e., once in T you remain in T
- Notation: for a node v in tree T , let $T(v)$ be the path from v to 0 along T .

Adding a node to the tree (continued)


- Consider any node u such that:
 - a) It has a neighbor v in T , and $(u, v)T(v)$ is allowed at u .
(notation: v is the first node in $T(v)$)
 - b) Of all the paths P in \mathcal{P}^u that are “consistent with T ” path $(u, v)T(v)$ (i.e. directly into T) is the highest ranked of these paths.
- If such u is found, add u to T
 - Note: such u may not exist.
 - In this case, we are “stuck”, and T does not become spanning.

Example: build the tree



- First: $T = \{0, x\}$ (x satisfies both (a) and (b) above)
- Note that, because of this, now ALL nodes u always have at least one path consistent with T, i.e., $(u, x, 0)$



Example: build the tree



- Next candidates:

- Because the tree is still small, all paths in all nodes are consistent with T.
- E.g., 1 has three consistent paths: $(1\ 3\ \underline{0})$, $(1\ \underline{0})$, $(1\ \underline{x}\ \underline{0})$
- However, in only node 3, the highest ranked one, i.e. $(3\ 0)$, has its next hop in T
- We thus add 3 to the tree.



Example continued ...



- Next candidates:

- Again, because the tree is still small, all paths in every node (not in T , i.e. in nodes 1, 2, 4, and 5) are consistent with T .
- E.g., 1 has three consistent paths: $(1\ \underline{3}\ \underline{0})$, $(1\ \underline{0})$, $(1\ \underline{x}\ \underline{0})$
- However, of these nodes, only node 1, the highest ranked one, i.e. $(1\ 3\ 0)$, has its next hop in T
- We thus add 1 to the tree.



Example continued ...



- Next candidates:

- All nodes not in T have a consistent path, but, not all paths in all nodes not in T (in nodes 2, 4, and 5) are consistent with T. E.g., (5 2 1 0) and (2 1 0) are not consistent since 1 has a tree path (1 3 0).
- Consistent paths: (2 0) (2 x 0) (4 2 0) (4 3 0) (4 x 0) (5 x 0)
- Of these nodes, 2 and 5 have their highest ranking path with a next hop in T. We can add either 2 or 5 to T.
- We arbitrarily choose to add 5 to the tree.



Example continued ...



- Next candidates:



- All nodes not in T have a consistent path, but, not all paths in all nodes not in T (in nodes 2, 4) are consistent with T. E.g., (2 1 0) is not consistent since 1 is taking the path (1 3 0).
- Consistent paths: (2 0) (2 x 0) (4 2 0) (4 3 0) (4 x 0)
- Of these nodes, 2 has its highest ranking path with a next hop in T.
- We add 2 to the tree.

Example continued ...




- Next candidates:
 - a) Only 4 remains. It has three consistent paths
(4 **2** 0) (4 **3** 0) (4 **x** 0)
 - b) Its highest ranked consistent path has its next hop in T.
 - c) We add 4 to the tree

Example continued ...



- T is a spanning tree
- T is stable (every step in construction yields a stable tree) (prove it!)
- Hence, there is a solution (T itself!)

What if we get stuck!




- Initially, $T = \{0, x\}$ (and all paths are consistent)
- Note, we can't add a single node more!
- We are stuck!
- All nodes prefer consistent paths not direct to x or 0 .

Stuck implies a dispute wheel . . .

- First, I am going to assume that in our SPP instance, if there is a path $(u_0, u_1, u_2, \dots, u_n)$ allowed at u_0 , then
 - a) path (u_1, u_2, \dots, u_n) is allowed at u_1 , and
 - b) each edge (u_i, u_{i+1}) , where $0 \leq i < n$, exists in graph G .
- Note that otherwise, path could be removed from the list in u_0 and it would not affect the converge behavior of the system (the path would never be taken)

Stuck implies a dispute wheel (contd)

- We will build one “spoke” of the wheel



- u_0, u_1, \dots, u_k are nodes not on the tree T .
- $\lambda^{u_0}(Q_0) < \lambda^{u_0}(R_0 Q_1)$

Stuck implies a dispute wheel (contd)

- Due to x , every u_0 not in T has at least one neighbor v_0 in T such that $(u_0, v_0)T(v_0)$ is allowed at u_0 .
 - a) Let $Q_0 = (u_0, v_0)T(v_0)$
- u_0 cannot be added to T . This implies that its highest ranking path P (consistent with T) has as next hop a node w that is not in T .
- P can be written as
 - a) $P : (u_0, w, \dots, u_1, v_1)T(v_1)$
 - b) where u_1 is the last node in P not in T ($u_1 = w$ is possible), and v_1 is in T
- Let $R_0 = (u_0, w, \dots, u_1)$, $Q_1 = (u_1, v_1)T(v_1)$, $P = R_0Q_1$
- From previous slide, $R_0Q_1 \in \mathcal{P}^{u_0}$ implies $Q_1 \in \mathcal{P}^{u_1}$

Stuck implies a dispute wheel (contd)

- Because P is the highest ranking consistent path at u_0 , we have
 - $\lambda(Q_0) < \lambda(P) = \lambda(R_0Q_1)$
- This completes one section of the dispute wheel
- Repeat the argument with u_1
 - note that
 - u_1 is not in T
 - Q_1 is permitted at u_1 , and its next hop (v_1) is in T ,
 - these are the same requirements we had on u_0 .
- As we repeat the argument the wheel gets built.
- Since the number of nodes is finite, we eventually end up at u_0 again.

No dispute wheel → single solution

- If there is no dispute wheel, then there is a single solution.
- PLEASE READ THIS ON YOUR OWN 😊

No dispute wheel \rightarrow Convergence

- Let α be a sequence of states such that
 - a) For each i , $\alpha(i+1)$ is obtained from $\alpha(i)$ by executing one node
 - b) For each i , $\alpha(i) \neq \alpha(i+1)$
- Let C be a cycle in α .
- A node u is *changing* in C if in two states in C it has different paths. Otherwise, u is *fixed*.
- Let $F(C)$ be the set of fixed nodes in C
- If $u \notin F(C)$, let $\text{values}(C, u)$ be the set of paths that u has taken in cycle C .

No dispute wheel \rightarrow Convergence

First, a Lemma:

- Let
 - a) u be a node that is not fixed in C ,
 - b) P be a path that is taken by u in C , (u changed its path from another path to P),
 - c) and v is the first “fixed node” of P .
- Then, each non-fixed node w , $w \in P[u, v]$, takes the path $P[w, 0]$ in some state in C .
- In particular, v stores $P[v, 0]$ throughout C .
- **READ THE PROOF OF THE LEMMA ON YOUR OWN**

No dispute wheel \rightarrow Convergence

Theorem: If there is a cycle C in an execution, then there exists a dispute wheel (no dispute wheel implies no cycles, and hence, convergence)

- Let u_0 be a node that
 - a) is not fixed in C , and
 - b) at some point in C , u_0 takes a path whose next hop w_0 is in $F(C)$, i.e., u_0 takes a “direct” path to $F(C)$.
 - c) By the Lemma, u_0 must exist. Why?
- Let Q_0 be the direct path of u_0 , i.e.,
 - a) $Q_0 = (u_0, w_0)Q_0'$, where $w_0 \in F(C)$, $Q_0' = \text{path of } w_0$.
- Q_0 is unique and the lowest ranked in $\text{values}(C, u_0)$. Why?


No dispute wheel \rightarrow Convergence . . .

- Let Z be the highest ranked path in $\text{values}(C, u_0)$
- From the lemma, Z consists of a sequence of nodes that are not fixed, followed by a sequence of nodes that are fixed
 - $Z = (u_0, \dots, u_1, w_1)Z'$ where:
 - u_0, \dots, u_1 are not fixed,
 - w_1 is fixed, and Z' is the path taken by w_1 .
- Let $R_0 = (u_0, \dots, u_1)$, and $Q_1 = (u_1, w_1)Z'$,
- We now have a spoke on the wheel
 - $\lambda(Q_0) < \lambda(Z) = \lambda(R_0 Q_1)$
- Repeat for the next spoke using u_1 and Q_1 .

Dispute Digraph


- The “nodes” in the graph correspond to paths in the SPP instance.
- It can be shown that a **cycle in the dispute graph** is the same as a **dispute wheel**.
- Hence, no dispute cycle iff no dispute wheel.
- There are two types of arcs (or edges) in the directed dispute graph:
 - a) Conflict arcs
 - b) Transmission arcs.

Dispute Arc




Note: it is also possible that $(u\ v)Q$ is not allowed at u

Gives the *dispute arc*




Transmission arc




Gives the *transmission arc*

$P \dashrightarrow (u,v)P$



Dispute Digraph Example



NAUGHTY GADGET



Dispute Digraph Example cont.



- Consider any node (say 2), think of it as u
- Take any path in u , say, 2 1 0 (**this is $(u v)P$**)
- The next node in the path (i.e. 1) is v
- If path 1 0 is in node 1 (it should be), then there is a **transmission arc** from 1 0 to 2 1 0.
- If there is a path in 1 with higher rank than 1 0, (yes there is, its 1 3 0), and 2 1 3 0 is ranked lower at 2 than 2 1 0 (it is lower ranked, since it is not even allowed at 2) then there is a **conflict arc** between 1 3 0 and 2 1 0.


NAUGHTY GADGET

An Application

c is a cost function on edges in SPP.

c is *coherent* if all cycles have positive cost (> 0).

An SPP specification is *consistent with c* if



consistent with coherent c

acyclic dispute digraph

will always converge

Dynamic execution model

- I WILL NOT COVER IT, AND IT WILL NOT BE IN THE EXAM
- However, feel free to read it for the fun of it 😊

Dynamic Execution Model

- For our purposes (to make our life easier) we will consider a “shared memory model” as opposed to a message passing model.
- Assume that each node can “read” the state of its neighbor (i.e. if $(u,v) \in E$, then u can read $\pi(v)$ and v can read $\pi(u)$).
- Execution is simple
 - a) Do forever:
 - If there is a node u such that $\pi(u) \neq \text{best}(\pi, u)$, then
 - a) $\pi(u) := \text{best}(\pi, u)$
 - b) We assume a “fair” execution that does not ignore some nodes.


A Dynamic Solution

- Extend SPP with a **history** attribute of each route
 - a) A route's history contains **a path in the dispute digraph** that “explains” how the route was obtained,
- Thus, if a route history contains a dispute cycle then a policy dispute was realized during the execution,
 - a) i.e. **the dispute graph has this cycle.**
- If a route's history contains a cycle, then suppress it (we will see how later)
- Nodes read their neighbor's current path and the associated path history of the path.

Updating the history

- If your new path is better than your previous path (higher ranked)
 - a) Add your NEW path (with a + sign) to the history of your NEW neighbor.
- If your new path is worse (lower ranked) than your previous path
 - a) Add your OLD path (with a – sign) to the (updated) history of your OLD neighbor
 - Note: you choose a new lower ranked path because the old path you had is no longer available.
 - Thus, the old neighbor changed paths (and of course also its history)

BAD GADGET



path

event history for path

1 : (1 0)	e
2 : (2 0)	e
3 : (3 4 2 0)	e
4 : (4 2 0)	e

execute 2

1 : (1 0)	e
2 : (2 1 0)	(+ 2 1 0)
3 : (3 4 2 0)	e
4 : (4 2 0)	e

execute 4

1 : (1 0)	e
2 : (2 1 0)	(+ 2 1 0)
3 : (3 4 2 0)	e
4 : ()	(- 4 2 0) (+ 2 1 0)

TIME



path

event history for path

execute 3

1 : (1 0)	e
2 : (2 1 0)	(+ 2 1 0)
3 : (3 0)	(- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
4 : ()	(- 4 2 0) (+ 2 1 0)

execute 1

1 : (1 3 0)	(+ 1 3 0) (- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
2 : (2 1 0)	(+ 2 1 0)
3 : (3 0)	(- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
4 : ()	(- 4 2 0) (+ 2 1 0)

execute 2

1 : (1 3 0)	(+ 1 3 0) (- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
2 : (2 0)	(- 2 1 0) (+ 1 3 0) (- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
3 : (3 0)	(- 3 4 2 0) (- 4 2 0) (+ 2 1 0)
4 : ()	(- 4 2 0) (+ 2 1 0)

TIME

A CYCLE!

What's going on ?

**Dynamic cycles of event history
correspond exactly to static cycles
in the dispute digraph.**